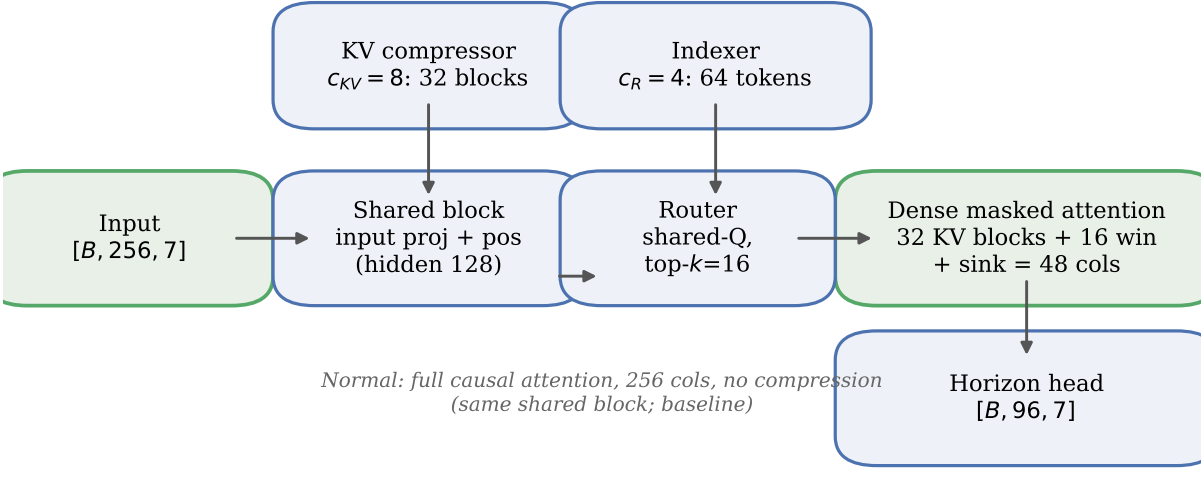
KV compressor
$c_{KV} = 8$: 32 blocks
Indexer
$c_R = 4$: 64 tokens
Input
$[B, 256, 7]$
Shared block
input proj + pos
(hidden 128)
Router
shared-Q,
top-k=16
Dense masked attention
32 KV blocks + 16 win
+ sink = 48 cols
Horizon head
$[B, 96, 7]$
Normal: full causal attention, 256 cols, no compression
(same shared block; baseline)