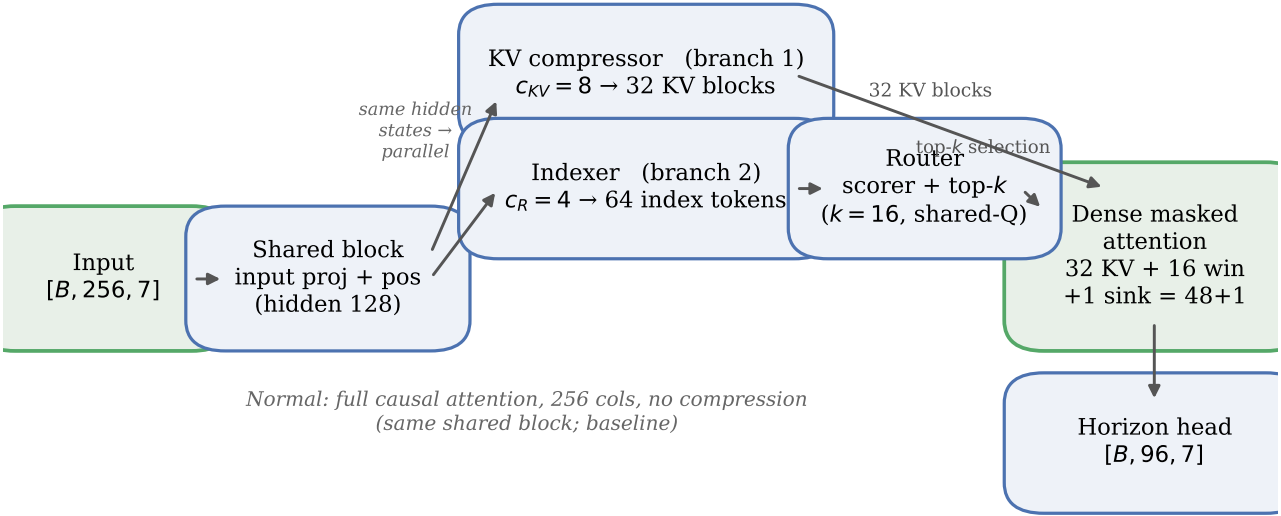
KV compressor (branch 1)
$c_{KV} = 8 \rightarrow 32$ KV blocks
same hidden states → parallel
Indexer (branch 2)
$c_R = 4 \rightarrow 64$ index tokens
32 KV blocks
top-$k$ selection
Router
scorer + top-$k$
($k = 16$, shared-Q)
Input
$[B, 256, 7]$
Shared block
input proj + pos
(hidden 128)
Dense masked attention
32 KV + 16 win
+1 sink = 48+1
Horizon head
$[B, 96, 7]$
Normal: full causal attention, 256 cols, no compression
(same shared block; baseline)