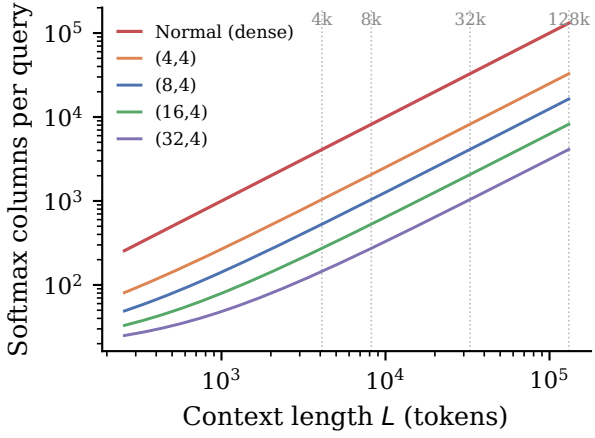
Softmax columns per query
Normal (dense)
(4,4)
(8,4)
(16,4)
(32,4)
4k
8k
32k
128k
Context length L (tokens)

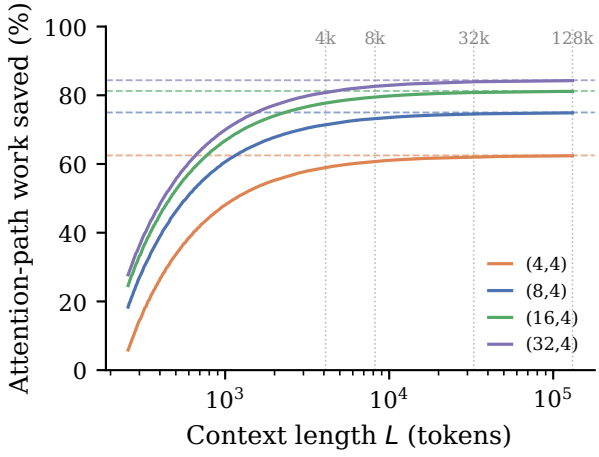
Attention-path work saved (%)
4k
8k
32k
128k
(4,4)
(8,4)
(16,4)
(32,4)
Context length L (tokens)

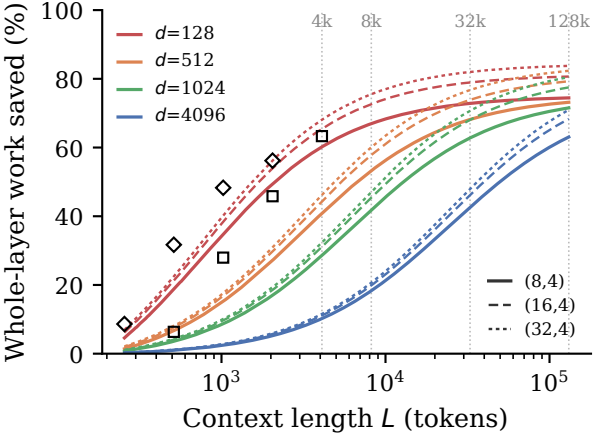
Whole-layer work saved (%)
d=128
d=512
d=1024
d=4096
4k
8k
32k
128k
(8,4)
(16,4)
(32,4)
Context length L (tokens)

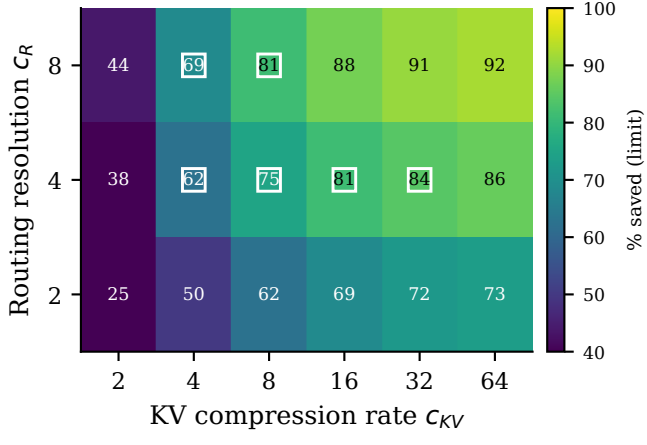
Routing resolution c_R
KV compression rate c_{KV}
% saved (limit)
44
69
81
88
91
92
38
62
75
81
84
86
25
50
62
69
72
73