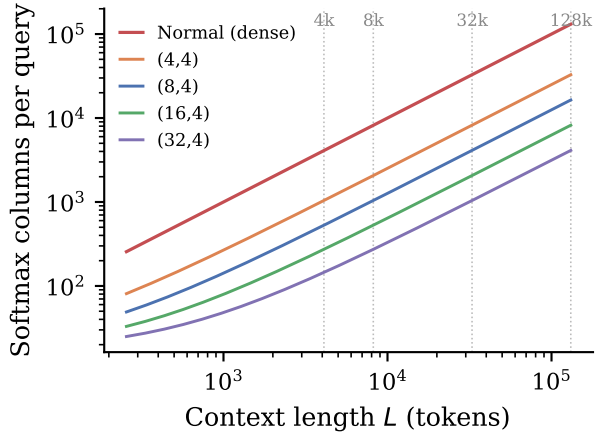
Softmax columns per query
Context length L (tokens)
Normal (dense)
(4,4)
(8,4)
(16,4)
(32,4)
4k
8k
32k
128k

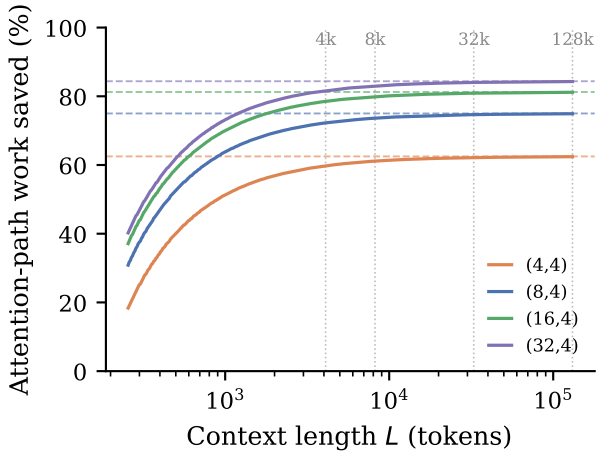
Attention-path work saved (%)
Context length L (tokens)
(4,4)
(8,4)
(16,4)
(32,4)
4k
8k
32k
128k

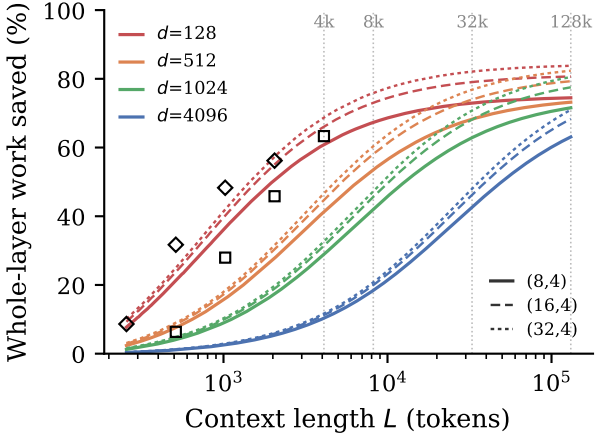
Whole-layer work saved (%)
Context length L (tokens)
d=128
d=512
d=1024
d=4096
(8,4)
(16,4)
(32,4)
4k
8k
32k
128k

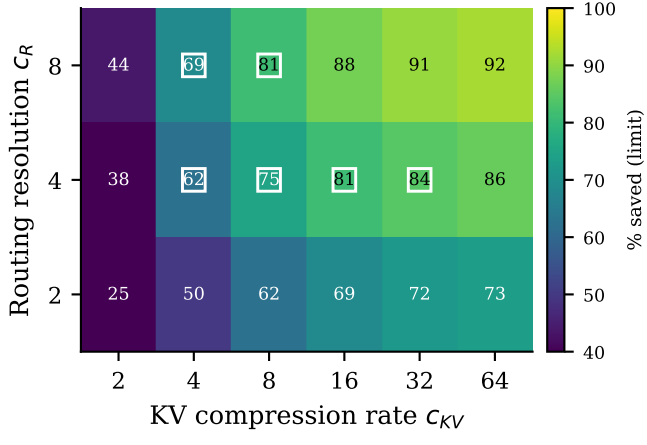
Routing resolution c_R
KV compression rate c_{KV}
% saved (limit)
44 69 81 88 91 92
38 62 75 81 84 86
25 50 62 69 72 73