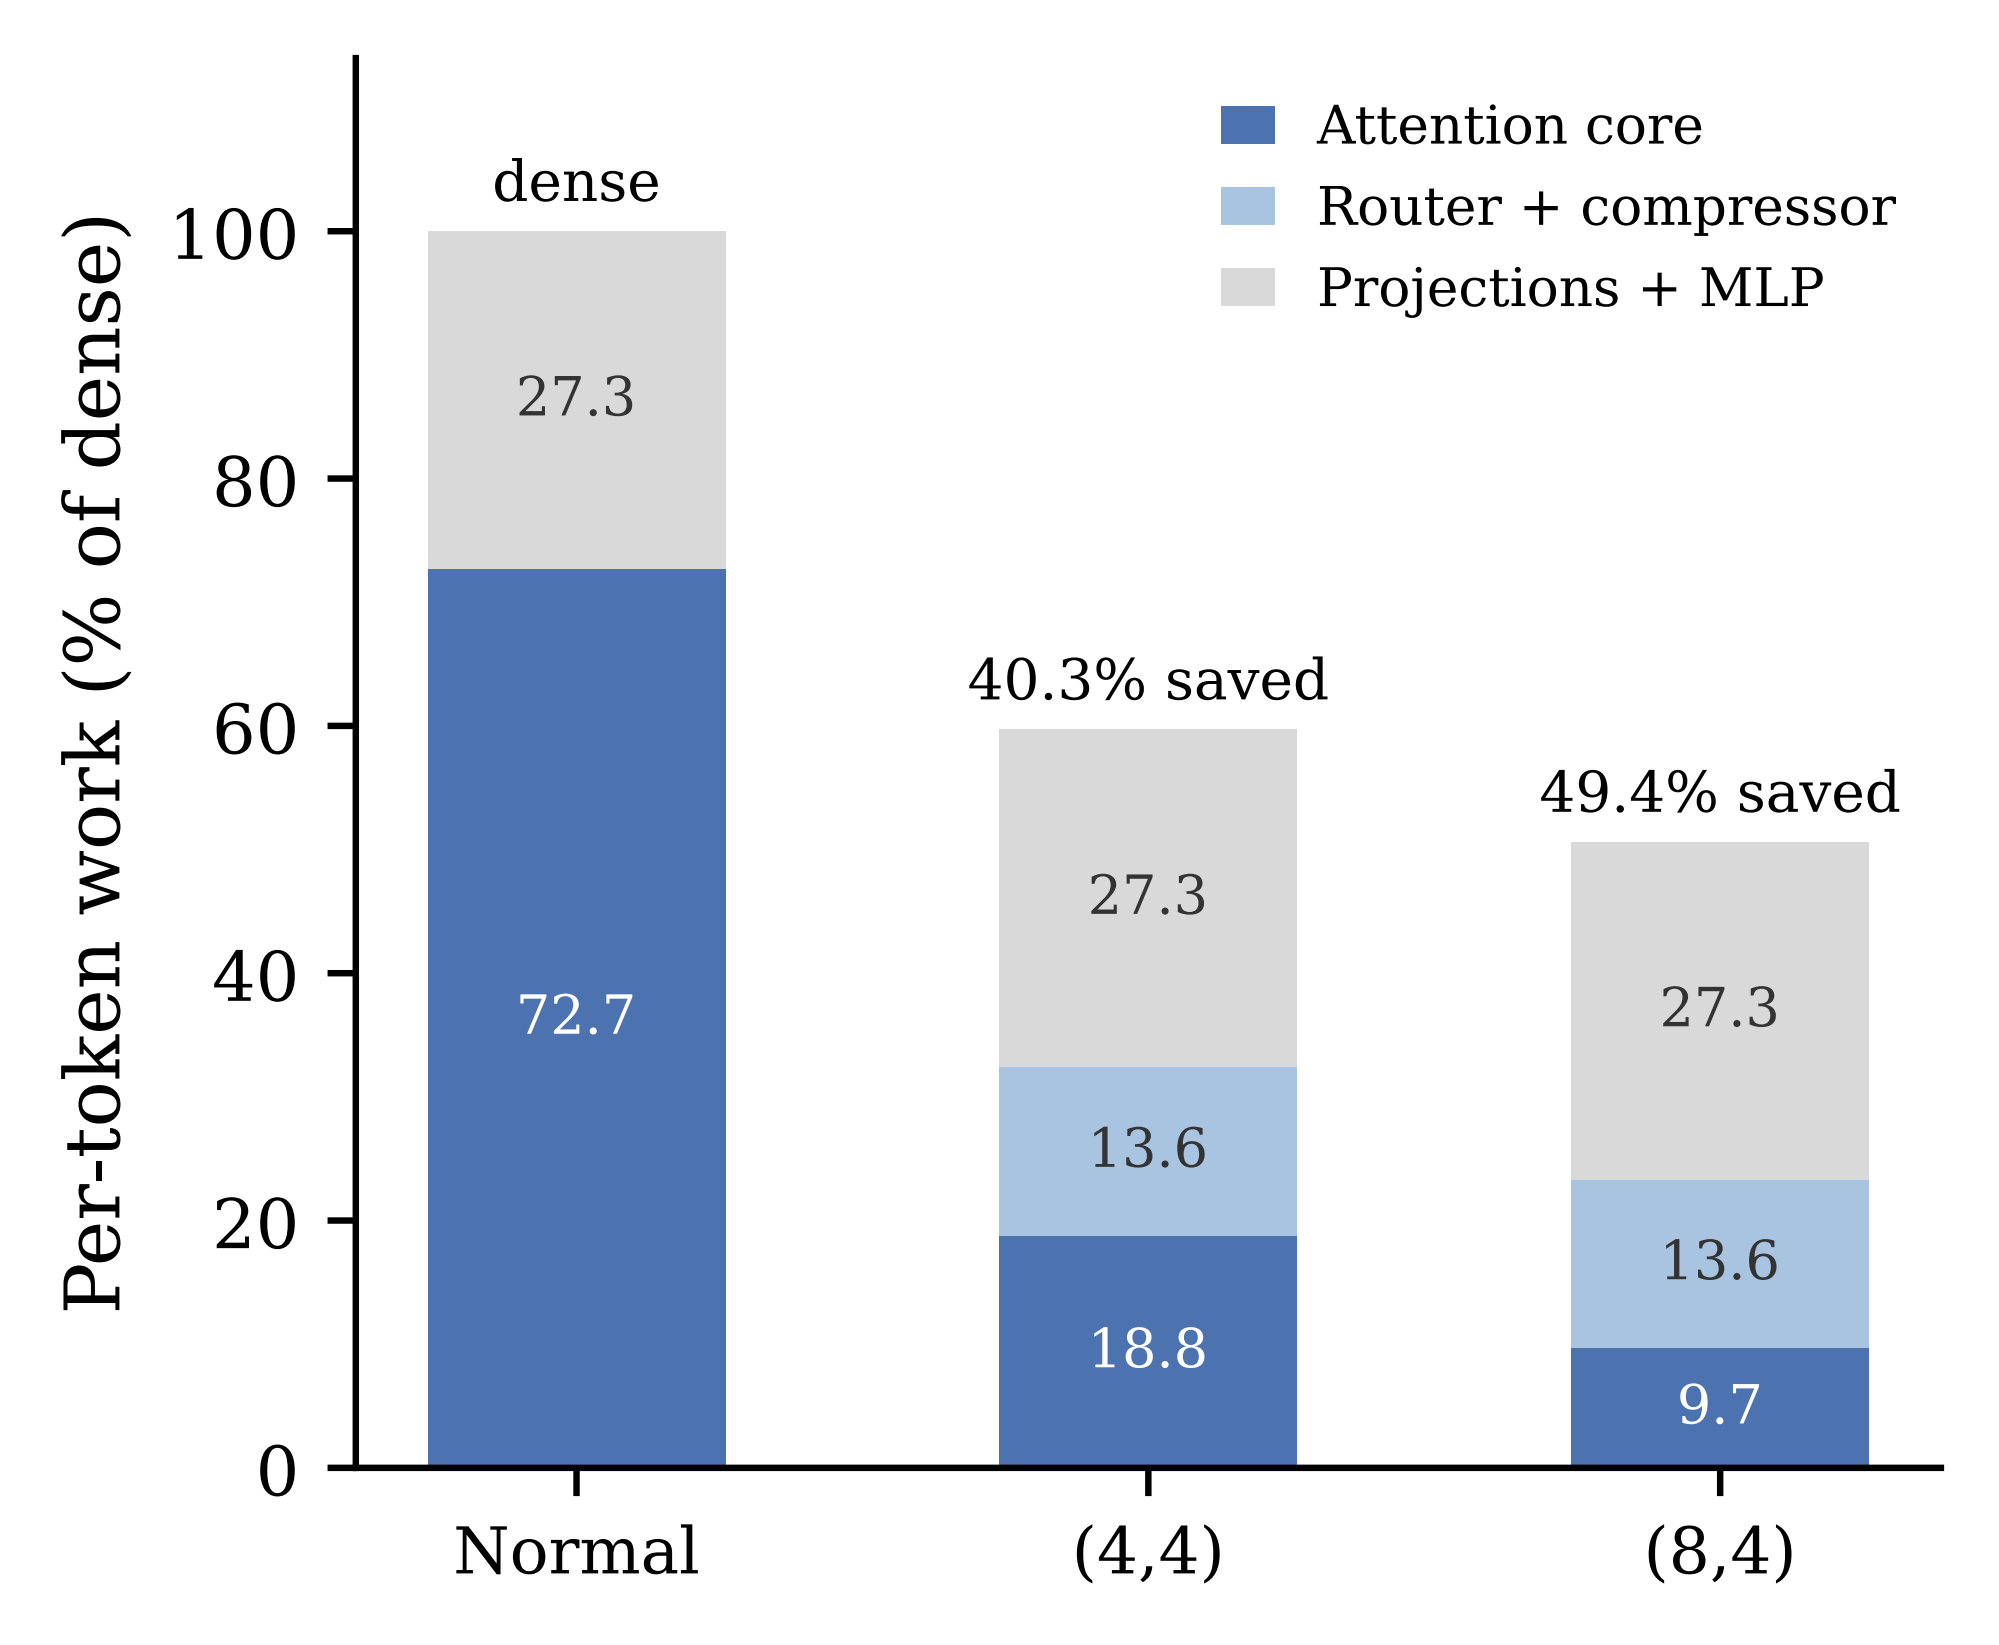
Attention core
Router + compressor
Projections + MLP
Per-token work (% of dense)
100
80
60
40
20
0
dense
27.3
72.7
40.3% saved
27.3
13.6
18.8
49.4% saved
27.3
13.6
9.7
Normal
(4,4)
(8,4)